

CRCLM_ml_0_1_2

May 9, 2023

```
[1]: import sys
    ## import all the packages needed
    import pandas as pd
    import numpy as np
    import matplotlib.pyplot as plt
    import sklearn
    import seaborn as sns
```

```
[2]: ## read NHANES dataset
    df = pd.read_csv('/Users/zhiiyi/Desktop/for Yupei/CRCLM_ml.csv')
```

```
[3]: ## find all the distinct values of DIQ170
    df.y.value_counts()
```

```
[3]: 0    23100
    1     5528
    Name: y, dtype: int64
```

```
[4]: ## data preparation
    # exclude null values and NA
    df = df[df.y.notnull()]
    # check DIQ170
    df.y.describe()
```

```
[4]: count    28628.000000
    mean         0.193098
    std          0.394736
    min          0.000000
    25%          0.000000
    50%          0.000000
    75%          0.000000
    max          1.000000
    Name: y, dtype: float64
```

```
[5]: # exclude non-numeric values
    d = df.select_dtypes(['number'])
```

```
# exclude columns that have over 50% NaN
d = d.dropna(thresh = 0.5*len(d), axis =1)

print(len(d.columns), 'columns are left')
```

21 columns are left

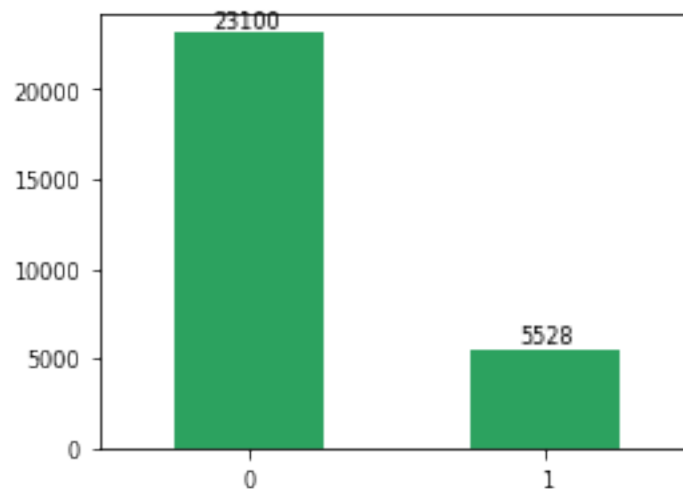
```
[6]: # transform the coding of target variable
d['y'] = d.apply(lambda x: 1 if x.y == 1 else 0, axis='columns')
d.y.value_counts()
```

```
[6]: 0    23100
      1     5528
      Name: y, dtype: int64
```

```
[7]: ## plot the distribution of values of response variable
vals = d.y.value_counts()

plt.figure(figsize=(4,3))
plt.rc('font', size=8)

ax = vals.plot.bar(rot=0, color='#2ca25f')
for i in range(len(vals)):
    ax.annotate(vals[i], xy=[vals.index[i], vals[i]], ha='center', va='bottom')
```



```
[8]: # replace NA with most frequent values
from sklearn.impute import SimpleImputer
sim_imp = SimpleImputer(strategy='most_frequent')

## avoid COD
```

```
d = pd.DataFrame(sim_imp.fit_transform(d), columns=d.columns).
↳drop(columns='COD')
```

```
## show the complete dataset
d
```

```
[8]:
```

	y	GENDER	RACE	AGE	SIZE	MARRY	INCOME	SITE	GRADE	KIND	T	N	\
0	1	2	2	3	3	1	3	1	9	1	3	3	
1	0	2	1	3	1	2	3	1	2	1	2	1	
2	0	1	1	3	1	3	3	2	2	1	2	2	
3	0	2	1	3	3	2	3	2	2	1	2	1	
4	0	1	3	3	1	3	3	1	9	1	2	1	
...
28623	1	2	2	3	1	3	1	2	2	2	2	1	
28624	0	1	1	2	1	1	1	1	9	1	3	3	
28625	1	1	2	2	1	2	2	1	9	1	3	1	
28626	1	2	1	3	2	3	2	2	2	1	2	1	
28627	1	1	2	3	3	2	2	2	9	1	3	3	
		SURGERY	RX	RADIATE	CHEM	CEA	BONE	BRAIN	LUNG				
0		0	0	0	0	1	0	0	1				
1		1	1	0	1	2	0	0	0				
2		1	0	0	0	0	0	0	0				
3		0	0	0	1	1	0	0	0				
4		0	0	0	0	1	0	0	1				
...				
28623		1	0	0	1	2	0	0	0				
28624		0	0	0	1	1	0	0	1				
28625		0	0	0	1	1	0	0	0				
28626		1	0	0	0	1	0	0	1				
28627		0	0	0	0	1	0	0	1				

[28628 rows x 20 columns]

```
[9]: ## separate predictors and responses
X = d.loc[:, d.columns != 'y']
Y = d.y
print('X shape:', X.shape)
print('Y shape:', Y.shape)
```

X shape: (28628, 19)

Y shape: (28628,)

```
[10]: ## split the data into training dataset and testing dataset (8:2)
from sklearn.model_selection import train_test_split
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size = 0.2,
↳random_state=12)
```

```
[11]: ## Feature selection: XGBClassifier
from xgboost import XGBClassifier
from sklearn.metrics import classification_report, accuracy_score, \
    confusion_matrix

model = XGBClassifier()
model.fit(X_train, Y_train)
Y_pred = model.predict(X_test)

accuracy = accuracy_score(Y_test, Y_pred)
print("Accuracy: %.2f%%" % (accuracy * 100.0))

## confusion matrix
def confusion(Y_test, Y_pred):
    conf = pd.DataFrame(confusion_matrix(Y_test, Y_pred), index=['True[0]', \
    'True[1]'], columns=['Predict[0]', 'Predict[1]'])
    print('Confusion Matrix:')
    print(conf)
    return conf

confusion(Y_test, Y_pred)
```

Accuracy: 82.55%

Confusion Matrix:

	Predict[0]	Predict[1]
True[0]	4301	304
True[1]	695	426

```
[11]:      Predict[0] Predict[1]
True[0]      4301        304
True[1]      695        426
```

```
[12]: ## Oversampling with SMOTE
from imblearn.over_sampling import SMOTE
smote = SMOTE()
X_train_smote, Y_train_smote = smote.fit_resample(X_train, Y_train)
X_test_smote, Y_test_smote = smote.fit_resample(X_test, Y_test)

X_train_smote = pd.DataFrame(X_train_smote, columns=X.columns)
X_test_smote = pd.DataFrame(X_test_smote, columns=X.columns)

## show the dataset after SMOTE
print(Y_train_smote.value_counts())
print(Y_test_smote.value_counts())
```

0	18495
1	18495

```
Name: y, dtype: int64
0    4605
1    4605
Name: y, dtype: int64
```

```
[13]: # The classification result after SMOTE
model = XGBClassifier()
model.fit(X_train_smote, Y_train_smote)
Y_pred_smote = model.predict(X_test_smote)

accuracy = accuracy_score(Y_test_smote, Y_pred_smote)
print("Accuracy: %.2f%%" % (accuracy * 100.0))
conf = pd.DataFrame(confusion_matrix(Y_test_smote, Y_pred_smote),
    ↪ index=['True[0]', 'True[1]'], columns=['Predict[0]', 'Predict[1]'])
conf
```

Accuracy: 75.75%

```
[13]:
```

	Predict[0]	Predict[1]
True[0]	3616	989
True[1]	1244	3361

```
[14]: from xgboost import XGBClassifier
from matplotlib import pyplot


# fit model no training data
model = XGBClassifier()
model.fit(X_train_smote, Y_train_smote)

# Features selected by XGBoost
keys = list(model.get_booster().feature_names)
values = list(model.feature_importances_)

data = pd.DataFrame(data=values, index=keys, columns=["score"]).sort_values(by=
    ↪ "score", ascending=False)

# Top xxx features
xgbfs_ = data[:]

# Plot feature score
xgbfs_.sort_values(by='score').plot(kind='barh', figsize=(10, 8),
    ↪ color='#2b8cbe')
plt.rc('font', size=8)
```



```
[15]: xgbfs = xgbfs_.reset_index()
xgbfs.columns=['variable', 'score']

xgbfs['variable'] = xgbfs['variable'].apply(lambda x: x.upper())
```

```
[16]: ## all variables list
var_list = xgbfs.variable.tolist()
var_list.append('y')
print(var_list)
df_final = d.filter(var_list)
df_final
```

```
['CHEM', 'SURGERY', 'RX', 'RADIATE', 'GRADE', 'RACE', 'CEA', 'INCOME', 'LUNG',
'AGE', 'N', 'T', 'SIZE', 'GENDER', 'SITE', 'MARRY', 'BONE', 'KIND', 'BRAIN',
'y']
```

```
[16]:
```

	CHEM	SURGERY	RX	RADIATE	GRADE	RACE	CEA	INCOME	LUNG	AGE	N	T	\
0	0	0	0	0	9	2	1	3	1	3	3	3	
1	1	1	1	0	2	1	2	3	0	3	1	2	
2	0	1	0	0	2	1	0	3	0	3	2	2	
3	1	0	0	0	2	1	1	3	0	3	1	2	

4	0	0	0	0	9	3	1	3	1	3	1	2
...
28623	1	1	0	0	2	2	2	1	0	3	1	2
28624	1	0	0	0	9	1	1	1	1	2	3	3
28625	1	0	0	0	9	2	1	2	0	2	1	3
28626	0	1	0	0	2	1	1	2	1	3	1	2
28627	0	0	0	0	9	2	1	2	1	3	3	3

	SIZE	GENDER	SITE	MARRY	BONE	KIND	BRAIN	y
0	3	2	1	1	0	1	0	1
1	1	2	1	2	0	1	0	0
2	1	1	2	3	0	1	0	0
3	3	2	2	2	0	1	0	0
4	1	1	1	3	0	1	0	0
...
28623	1	2	2	3	0	2	0	1
28624	1	1	1	1	0	1	0	0
28625	1	1	1	2	0	1	0	1
28626	2	2	2	3	0	1	0	1
28627	3	1	2	2	0	1	0	1


[28628 rows x 20 columns]

```
[17]: ## draw correlation heatmap
ax = plt.subplots(figsize=(10,10))
plt.rc('font', size=9)

corr = df_final.corr()

sns.heatmap(corr, cmap="Blues")
```

[17]: <AxesSubplot:>



```
[18]: X_scale = df_final.loc[:, df_final.columns != 'y']
      Y = df_final.y
```

```
[19]: ## min-max scaling
      from sklearn.preprocessing import MinMaxScaler
      minmax = MinMaxScaler()
      X = pd.DataFrame(minmax.fit_transform(X_scale), columns=X_scale.columns)
      X
```

```
[19]:
```

	CHEM	SURGERY	RX	RADIATE	GRADE	RACE	CEA	INCOME	LUNG	AGE	N	\
0	0.0	0.0	0.0	0.0	1.000	0.5	0.5	1.0	0.5	1.0	1.0	
1	1.0	1.0	1.0	0.0	0.125	0.0	1.0	1.0	0.0	1.0	0.0	

2	0.0	1.0	0.0	0.0	0.125	0.0	0.0	1.0	0.0	1.0	0.5
3	1.0	0.0	0.0	0.0	0.125	0.0	0.5	1.0	0.0	1.0	0.0
4	0.0	0.0	0.0	0.0	1.000	1.0	0.5	1.0	0.5	1.0	0.0
...
28623	1.0	1.0	0.0	0.0	0.125	0.5	1.0	0.0	0.0	1.0	0.0
28624	1.0	0.0	0.0	0.0	1.000	0.0	0.5	0.0	0.5	0.5	1.0
28625	1.0	0.0	0.0	0.0	1.000	0.5	0.5	0.5	0.0	0.5	0.0
28626	0.0	1.0	0.0	0.0	0.125	0.0	0.5	0.5	0.5	1.0	0.0
28627	0.0	0.0	0.0	0.0	1.000	0.5	0.5	0.5	0.5	1.0	1.0

	T	SIZE	GENDER	SITE	MARRY	BONE	KIND	BRAIN
0	1.0	1.0	1.0	0.0	0.0	0.0	0.0	0.0
1	0.5	0.0	1.0	0.0	0.5	0.0	0.0	0.0
2	0.5	0.0	0.0	1.0	1.0	0.0	0.0	0.0
3	0.5	1.0	1.0	1.0	0.5	0.0	0.0	0.0
4	0.5	0.0	0.0	0.0	1.0	0.0	0.0	0.0
...
28623	0.5	0.0	1.0	1.0	1.0	0.0	1.0	0.0
28624	1.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
28625	1.0	0.0	0.0	0.0	0.5	0.0	0.0	0.0
28626	0.5	0.5	1.0	1.0	1.0	0.0	0.0	0.0
28627	1.0	1.0	0.0	1.0	0.5	0.0	0.0	0.0

[28628 rows x 19 columns]

```
[20]: ## Create an empty list containing the accuracy scores of the models
mscore=[]
```

```
[21]: from sklearn.model_selection import train_test_split
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, random_state=11)

X_train_smote, Y_train_smote = smote.fit_resample(X_train, Y_train)
X_test_smote, Y_test_smote = smote.fit_resample(X_test, Y_test)

X_train_smote = pd.DataFrame(X_train_smote, columns=X.columns)
X_test_smote = pd.DataFrame(X_test_smote, columns=X.columns)
```

```
[22]: ## Logistic Regression
from sklearn.linear_model import LogisticRegression
import time

clf = LogisticRegression(max_iter=100, solver='lbfgs', class_weight='balanced',
    random_state=11)

## time of running
tic = time.time()
clf.fit(X_train_smote, Y_train_smote)
```

```

toc = time.time()
print("Time: " + str(1000*(toc-tic)) + "ms")

clf_prediction_proba = clf.predict_proba(X_test)[: , 1]

## make predictions
Y_pred = clf.predict(X_test)
print('Accuracy Score:', clf.score(X_test, Y_test))
print('Prediction:', Y_pred)

## append the score to the list
mscore.append(['Logistic Regression', clf.score(X_test, Y_test)])

print(classification_report(Y_test, Y_pred))
confusion(Y_test, Y_pred)

```

Time: 144.3309783935547ms

Accuracy Score: 0.7510129942713427

Prediction: [0 0 0 ... 0 0 1]

	precision	recall	f1-score	support
0	0.91	0.76	0.83	5739
1	0.42	0.70	0.53	1418
accuracy			0.75	7157
macro avg	0.67	0.73	0.68	7157
weighted avg	0.82	0.75	0.77	7157

Confusion Matrix:

	Predict[0]	Predict[1]
True[0]	4377	1362
True[1]	420	998

```

[22]:
      Predict[0] Predict[1]
True[0]      4377      1362
True[1]      420       998

```

```

[23]: ## random forest model
from sklearn.ensemble import RandomForestClassifier
import time

rnd_clf = RandomForestClassifier(n_estimators=200, criterion='gini',
    ↪max_depth=5,
                                min_samples_split = 2, min_samples_leaf = 1,
    ↪random_state=11)

## time of running

```

```

tic = time.time()
rnd_clf.fit(X_train_smote, Y_train_smote)
toc = time.time()
print("Time: " + str(1000*(toc-tic)) + "ms")

rf_prediction_proba = rnd_clf.predict_proba(X_test)[:, 1]

## make predictions
Y_pred = rnd_clf.predict(X_test)
print('Accuracy Score:', rnd_clf.score(X_test, Y_test))
print('Prediction:', Y_pred)

## append the score to the list
mscore.append(['Random Forest', rnd_clf.score(X_test, Y_test)])

print(classification_report(Y_test, Y_pred))
confusion(Y_test, Y_pred)

## print importance tree random forest
print(pd.DataFrame({'Variable':X.columns,
                    'Importance':model.feature_importances_}).
      ↪sort_values('Importance', ascending=False))

```

Time: 3028.102159500122ms

Accuracy Score: 0.769037306133855

Prediction: [0 0 0 ... 0 0 1]

	precision	recall	f1-score	support
0	0.91	0.80	0.85	5739
1	0.44	0.66	0.53	1418
accuracy			0.77	7157
macro avg	0.67	0.73	0.69	7157
weighted avg	0.81	0.77	0.78	7157

Confusion Matrix:

	Predict[0]	Predict[1]
True[0]	4563	1176
True[1]	477	941

	Variable	Importance
14	SITE	0.602844
11	T	0.114928
12	SIZE	0.030765
13	GENDER	0.029534
7	INCOME	0.020103
1	SURGERY	0.017172
15	MARRY	0.017033

5	RACE	0.016895
18	BRAIN	0.016194
2	RX	0.015621
10	N	0.015263
9	AGE	0.014503
3	RADIATE	0.014333
0	CHEM	0.013556
6	CEA	0.013118
4	GRADE	0.012994
16	BONE	0.012251
8	LUNG	0.011788
17	KIND	0.011105

```
[24]: ## gradient boosting
from sklearn.ensemble import GradientBoostingClassifier
import time

gbc = GradientBoostingClassifier(learning_rate=0.1, n_estimators=10,
    random_state=11)

## time of running
tic = time.time()
gbc.fit(X_train_smote, Y_train_smote)
toc = time.time()
print("Time: " + str(1000*(toc-tic)) + "ms")

gbc_prediction_proba = gbc.predict_proba(X_test)[: , 1]

## make predictions
Y_pred = gbc.predict(X_test)
print('Accuracy Score:', gbc.score(X_test, Y_test))
print('Prediction:', Y_pred)

## append the score to the list
mscore.append(['Gradient Boosting', gbc.score(X_test, Y_test)])

print(classification_report(Y_test, Y_pred))
confusion(Y_test, Y_pred)
```

Time: 445.06025314331055ms

Accuracy Score: 0.8006147827301943

Prediction: [0 0 0 ... 0 0 0]

	precision	recall	f1-score	support
0	0.89	0.85	0.87	5739
1	0.50	0.59	0.54	1418
accuracy			0.80	7157

macro avg	0.70	0.72	0.71	7157
weighted avg	0.82	0.80	0.81	7157

Confusion Matrix:

	Predict[0]	Predict[1]
True[0]	4888	851
True[1]	576	842

```
[24]:      Predict[0] Predict[1]
True[0]      4888      851
True[1]      576      842
```

```
[25]: ## Adaboost Classifier
from sklearn.ensemble import AdaBoostClassifier
import time

ada = AdaBoostClassifier(learning_rate=0.01, n_estimators=30, random_state=11)

tic = time.time()
ada.fit(X_train_smote, Y_train_smote)
toc = time.time()
print("Time: " + str(1000*(toc-tic)) + "ms")

ada_prediction_proba = ada.predict_proba(X_test)[: , 1]

Y_pred = ada.predict(X_test)
print('Accuracy Score:', ada.score(X_test, Y_test))
print('Prediction:', Y_pred)

mscore.append(['Adaptive Boosting', ada.score(X_test, Y_test)])

# from sklearn.metrics import classification_report
print(classification_report(Y_test, Y_pred))
confusion(Y_test, Y_pred)
```

Time: 902.2519588470459ms

Accuracy Score: 0.797820315774766

Prediction: [0 0 0 ... 0 0 0]

	precision	recall	f1-score	support
0	0.90	0.85	0.87	5739
1	0.49	0.60	0.54	1418
accuracy			0.80	7157
macro avg	0.69	0.72	0.71	7157
weighted avg	0.82	0.80	0.80	7157

Confusion Matrix:

	Predict[0]	Predict[1]
True[0]	4861	878
True[1]	569	849

```
[25]:
```

	Predict[0]	Predict[1]
True[0]	4861	878
True[1]	569	849

```
[26]: ## Extreme Gradient Boosting
from xgboost import XGBClassifier
import time

xgbc = XGBClassifier(eta=0.01, max_depth=3, random_state=11)

tic = time.time()
xgbc.fit(X_train_smote, Y_train_smote)
toc = time.time()
print("Time: " + str(1000*(toc-tic)) + "ms")

xgb_prediction_proba = xgbc.predict_proba(X_test)[: , 1]

Y_pred = xgbc.predict(X_test)
print('Accuracy Score:', xgbc.score(X_test, Y_test))
print('Prediction:', Y_pred)

mscore.append(['XGBoost', xgbc.score(X_test, Y_test)])

# from sklearn.metrics import classification_report
print(classification_report(Y_test, Y_pred))
confusion(Y_test, Y_pred)
```

Time: 2475.6009578704834ms

Accuracy Score: 0.7989381025569373

Prediction: [0 0 0 ... 0 0 0]

	precision	recall	f1-score	support
0	0.89	0.85	0.87	5739
1	0.49	0.60	0.54	1418
accuracy			0.80	7157
macro avg	0.69	0.72	0.71	7157
weighted avg	0.82	0.80	0.81	7157

Confusion Matrix:

	Predict[0]	Predict[1]
True[0]	4872	867
True[1]	572	846

```
[26]:          Predict[0] Predict[1]
True[0]          4872         867
True[1]          572         846
```

```
[27]: ## Support Vector Machine
from sklearn.svm import SVC
import time

svm_clf = SVC(kernel='sigmoid', gamma='auto', random_state=11, probability=True)

## time of running
tic = time.time()
svm_clf.fit(X_train_smote, Y_train_smote)
toc = time.time()
print("Time: " + str(1000*(toc-tic)) + "ms")

svm_prediction_proba = svm_clf.predict_proba(X_test)[:, 1]

Y_pred = svm_clf.predict(X_test)
print('Accuracy Score:', svm_clf.score(X_test, Y_test))
print('Prediction:', Y_pred)

mscore.append(['SVM', svm_clf.score(X_test, Y_test)])

# from sklearn.metrics import classification_report
print(classification_report(Y_test, Y_pred))
confusion(Y_test, Y_pred)
```

Time: 725942.412853241ms

Accuracy Score: 0.6797540869079223

Prediction: [0 0 1 ... 1 1 0]

	precision	recall	f1-score	support
0	0.89	0.68	0.77	5739
1	0.34	0.67	0.45	1418
accuracy			0.68	7157
macro avg	0.62	0.68	0.61	7157
weighted avg	0.78	0.68	0.71	7157

Confusion Matrix:

	Predict[0]	Predict[1]
True[0]	3914	1825
True[1]	467	951

```
[27]:          Predict[0] Predict[1]
True[0]          3914         1825
```

True[1] 467 951

```
[28]: ## Bagging K-Nearest Neighbors
from sklearn.ensemble import BaggingClassifier
from sklearn.neighbors import KNeighborsClassifier
import time

bagging = BaggingClassifier(base_estimator= KNeighborsClassifier(), max_samples=
    ↳ 0.5, max_features = 0.5,
                                bootstrap = False, bootstrap_features = False,
    ↳ random_state=11)

## time of running
tic = time.time()
bagging.fit(X_train_smote, Y_train_smote)
toc = time.time()
print("Time: " + str(1000*(toc-tic)) + "ms")

kn_prediction_proba = bagging.predict_proba(X_test)[: , 1]

## make predictions
Y_pred = bagging.predict(X_test)
print('Accuracy Score:', bagging.score(X_test, Y_test))
print('Prediction:', Y_pred)

bg_score = bagging.score(X_test, Y_test)
bagging.score(X_test, Y_test)

## append the score to the list
mscore.append(['Bagging_KNeighbors', bagging.score(X_test, Y_test)])
```

Time: 294.31796073913574ms
Accuracy Score: 0.8043873131200223
Prediction: [0 0 0 ... 0 0 1]

```
[29]: ## Bagging Decision Tree
from sklearn.tree import DecisionTreeClassifier
import time

bagging = BaggingClassifier(base_estimator= DecisionTreeClassifier(),
    ↳ max_samples = 0.5, max_features = 0.5,
                                bootstrap = False, bootstrap_features = False,
    ↳ random_state=11)

## time of running
tic = time.time()
bagging.fit(X_train_smote, Y_train_smote)
```



```

toc = time.time()
print("Time:" + str(1000*(toc-tic)) + "ms")

bdt_prediction_proba = bagging.predict_proba(X_test)[: , 1]

## make predictions
Y_pred = bagging.predict(X_test)
print('Accuracy Score:', bagging.score(X_test, Y_test))
print('Prediction:', Y_pred)

bg_dt_score = bagging.score(X_test, Y_test)
bagging.score(X_test, Y_test)

## append the score to the list
mscore.append(['Bagging Decision Tree based', bagging.score(X_test, Y_test)])

```

Time:373.25477600097656ms
Accuracy Score: 0.8129104373340785
Prediction: [0 0 0 ... 0 0 0]

```

[30]: ## MLP
from sklearn.neural_network import MLPClassifier
import time

mlp = MLPClassifier(hidden_layer_sizes=(100,100), solver='adam', shuffle=False,
    ↪tol = 0.0001, random_state=11)

## time of running
tic = time.time()
mlp.fit(X_train_smote, Y_train_smote)
toc = time.time()
print("Time:" + str(1000*(toc-tic)) + "ms")

mlp_prediction_proba = mlp.predict_proba(X_test)[: , 1]

## make predictions
Y_pred = mlp.predict(X_test)
print('Accuracy Score:', mlp.score(X_test, Y_test))
print('Prediction:', Y_pred)

## see the parameters ready to be adjusted
print("parameter: ", mlp.get_params())

## append the score to the list
mscore.append(['MLP', mlp.score(X_test, Y_test)])

```

Time:4644.575834274292ms

Accuracy Score: 0.4223836803129803

Prediction: [1 1 1 ... 1 0 1]

```
parameter: {'activation': 'relu', 'alpha': 0.0001, 'batch_size': 'auto',
'beta_1': 0.9, 'beta_2': 0.999, 'early_stopping': False, 'epsilon': 1e-08,
'hidden_layer_sizes': (100, 100), 'learning_rate': 'constant',
'learning_rate_init': 0.001, 'max_fun': 15000, 'max_iter': 200, 'momentum': 0.9,
'n_iter_no_change': 10, 'nesterovs_momentum': True, 'power_t': 0.5,
'random_state': 11, 'shuffle': False, 'solver': 'adam', 'tol': 0.0001,
'validation_fraction': 0.1, 'verbose': False, 'warm_start': False}
```

```
[31]: mscore.sort(key=lambda x: x[1], reverse=True)
mscore
```

```
[31]: [['Bagging_Decision Tree based', 0.8129104373340785],
['Bagging_KNeighbors', 0.8043873131200223],
['Gradient Boosting', 0.8006147827301943],
['XGBoost', 0.7989381025569373],
['Adaptive Boosting', 0.797820315774766],
['Random Forest', 0.769037306133855],
['Logistic Regression', 0.7510129942713427],
['SVM', 0.6797540869079223],
['MLP', 0.4223836803129803]]
```

```
[32]: model = list(i[0] for i in mscore)
score = list(round(i[1]*100,2) for i in mscore)

print('Accuracy Score: \n')
for m,s in zip(model, score):
    print(f'{m}: {s}%')

# creating horizontal bar
plt.barh(model, score, height = 0.5, color='#4B4E6D')

plt.xlabel("Accuracy Score (%)")
plt.ylabel("Model")
plt.title("Model Comparison regarding Accuracy")
plt.gca().invert_yaxis()
plt.rc('font', size=9)
plt.show()
```

Accuracy Score:

Bagging_Decision Tree based: 81.29%

Bagging_KNeighbors: 80.44%


Gradient Boosting: 80.06%

XGBoost: 79.89%

Adaptive Boosting: 79.78%

Random Forest: 76.9%

Logistic Regression: 75.1%
SVM: 67.98%
MLP: 42.24%



```
[33]: from sklearn.metrics import roc_auc_score, roc_curve
import matplotlib
import matplotlib.pyplot as plt
plt.figure(figsize=(9, 7))
matplotlib.rcParams.update({'font.size': 15})
plt.grid()

def roc_curve_and_score(y_test, pred_proba):
    fpr, tpr, _ = roc_curve(y_test.ravel(), pred_proba.ravel())
    roc_auc = roc_auc_score(y_test.ravel(), pred_proba.ravel())
    return fpr, tpr, roc_auc

## ROC curve
### Bagging Decision Tree
fpr, tpr, roc_auc = roc_curve_and_score(Y_test, bdt_prediction_proba)
plt.plot(fpr, tpr, color='#ff7f00', lw=2,
         label='Bagging Decision Tree={0:.3f}'.format(roc_auc))
### Random Forest
fpr, tpr, roc_auc = roc_curve_and_score(Y_test, rf_prediction_proba)
plt.plot(fpr, tpr, color='#984ea3', lw=2,
         label='Random Forest={0:.3f}'.format(roc_auc))
### Bagging KNeighbors
fpr, tpr, roc_auc = roc_curve_and_score(Y_test, kn_prediction_proba)
```


```

plt.plot(fpr, tpr, color='#4daf4a', lw=2,
         label='Bagging KNeighbors={0:.3f}'.format(roc_auc))
### Gradient Boosting
fpr, tpr, roc_auc = roc_curve_and_score(Y_test, gbc_prediction_proba)
plt.plot(fpr, tpr, color='#377eb8', lw=2,
         label='Gradient Boosting={0:.3f}'.format(roc_auc))
### AdaBoost
fpr, tpr, roc_auc = roc_curve_and_score(Y_test, ada_prediction_proba)
plt.plot(fpr, tpr, color='#00FFFF', lw=2,
         label='AdaBoost={0:.3f}'.format(roc_auc))
### XGBoost
fpr, tpr, roc_auc = roc_curve_and_score(Y_test, xgb_prediction_proba)
plt.plot(fpr, tpr, color='#E3CF57', lw=2,
         label='XGBoost={0:.3f}'.format(roc_auc))
### SVM
fpr, tpr, roc_auc = roc_curve_and_score(Y_test, svm_prediction_proba)
plt.plot(fpr, tpr, color='#808A87', lw=2,
         label='SVM={0:.3f}'.format(roc_auc))
### MLP
fpr, tpr, roc_auc = roc_curve_and_score(Y_test, mlp_prediction_proba)
plt.plot(fpr, tpr, color='#8A2BE2', lw=2,
         label='MLP={0:.3f}'.format(roc_auc))
### Logistic Regression
fpr, tpr, roc_auc = roc_curve_and_score(Y_test, clf_prediction_proba)
plt.plot(fpr, tpr, color='#e41a1c', lw=2, linestyle='--',
         label='Logistic Regression={0:.3f}'.format(roc_auc))

## reference line (diagonal)
plt.plot([0, 1], [0, 1], color='black', lw=2, linestyle='--')

## add legends, labels, and plot the ROC curves
plt.legend(loc="lower right", fontsize="12")
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('1 - Specificity (False Positive Rate)')
plt.ylabel('Sensitivity (True Positive Rate)')
plt.title("ROC curve")
plt.rc('font', size=9)
plt.show()

```



[]: